

In [2]: *# You can load and split the dataset using this code block:*

```
from sklearn.datasets import load_digits
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
digits= load_digits()
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression, Ridge, Lasso
from sklearn.metrics import confusion_matrix


from sklearn.model_selection import StratifiedKFold

x = digits.data
y = digits.target

x_train, x_test, y_train, y_test = train_test_split(digits.data, digits.target
, test_size=0.20, random_state=0)

# You can view the images using this code snippet:

plt.imshow(np.reshape(x_train[7], (8,8)), cmap=plt.cm.gray)
plt.show()
print(digits.data[100])
print(digits.target[100])
print(len(digits.target))
```



```
[ 0.  0.  0.  2. 13.  0.  0.  0.  0.  0.  0.  8. 15.  0.  0.
  0.  0.  0.  5. 16.  5.  2.  0.  0.  0.  0. 15. 12.  1. 16.
  4.  0.  0.  4. 16.  2.  9. 16.  8.  0.  0.  0. 10. 14. 16.
 16.  4.  0.  0.  0.  0.  0. 13.  8.  0.  0.  0.  0.  0.  0.
 13.  6.  0.  0.]
4
1797
```

```

In [91]: logreg_ridge = LogisticRegression() # instantiate

c_vals = np.linspace(0.1,100,10)

skf = StratifiedKFold(n_splits=5)
skf.get_n_splits(x, y)
scores_l2_train = []
scores_l2_test = []
scores_l2_c = []
scores_l2_split = []
split_count = 0

for c in c_vals:
    #     print("L2 Logistic Regression (Ridge)\n")
    split_count = 0
    logreg_ridge = LogisticRegression(C=c)
    for train_index, test_index in skf.split(x, y):
        #     print("TRAIN:", train_index, "TEST:", test_index)
        x_train, x_test = x[train_index], x[test_index]
        y_train, y_test = y[train_index], y[test_index]

        logreg_ridge.fit(x_train, y_train) # fit
        y_pred = logreg_ridge.predict(x_train) # predict
        acc_log = sum(y_pred == y_train)/len(y_train)*100
        #     print('Logistic Regression accuracy of train data:', str(round(acc_log,2)), '%')
        #     print(confusion_matrix(y_train, y_pred))
        scores_l2_train.append(acc_log)

        y_pred = logreg_ridge.predict(x_test) # predict
        acc_log = sum(y_pred == y_test)/len(y_test)*100
        #     print('Logistic Regression accuracy of test data:', str(round(acc_log,2)), '%')
        #     print(confusion_matrix(y_test, y_pred))

    #     print("\n")
    scores_l2_test.append(acc_log)
    scores_l2_c.append(c)
    split_count = split_count + 1
    scores_l2_split.append(split_count)

```

```
In [92]: data = {"train": scores_l2_train, "test": scores_l2_test, "c": scores_l2_c, "split": scores_l2_split}

# scores_l2 = pd.DataFrame([scores_l2_train, scores_l2_test, scores_l2_c])
scores_l2 = pd.DataFrame(data=data)

scores_l2[scores_l2.columns[:-1]]
# scores_l2

#     logreg_lasso = LogisticRegression(penalty='L1')
#     print("L1 Logistic Regression (Lasso)\n")
#     for train_index, test_index in skf.split(x, y):
#         #     print("TRAIN:", train_index, "TEST:", test_index)
#         x_train, x_test = x[train_index], x[test_index]
#         y_train, y_test = y[train_index], y[test_index]

#         logreg_lasso.fit(x_train, y_train) # fit
#         y_pred = logreg_lasso.predict(x_train) # predict
#         acc_log = sum(y_pred == y_train)/len(y_train)*100
#         print('Logistic Regression accuracy of train data:', str(round(acc_log,2)), '%')
#         #     print(confusion_matrix(y_train, y_pred))

#         y_pred = logreg_lasso.predict(x_test) # predict
#         acc_log = sum(y_pred == y_test)/len(y_test)*100
#         print('Logistic Regression accuracy of test data:', str(round(acc_log,2)), '%')
#         #     print(confusion_matrix(y_test, y_pred))

#     print("\n")
```

Out[92]:

	<b>train</b>	<b>test</b>	<b>split</b>	<b>c</b>
<b>0</b>	99.302163	92.032967	1	0.1
<b>1</b>	99.303136	88.674033	2	0.1
<b>2</b>	99.513213	96.100279	3	0.1
<b>3</b>	99.375000	96.078431	4	0.1
<b>4</b>	99.445215	89.295775	5	0.1
<b>5</b>	99.790649	91.758242	1	11.2
<b>6</b>	99.790941	87.292818	2	11.2
<b>7</b>	99.791377	94.150418	3	11.2
<b>8</b>	99.791667	95.238095	4	11.2
<b>9</b>	100.000000	90.422535	5	11.2
<b>10</b>	99.790649	91.758242	1	22.3
<b>11</b>	99.860627	87.569061	2	22.3
<b>12</b>	99.791377	93.593315	3	22.3
<b>13</b>	99.791667	95.238095	4	22.3
<b>14</b>	100.000000	89.859155	5	22.3
<b>15</b>	99.860433	91.758242	1	33.4
<b>16</b>	99.860627	87.292818	2	33.4
<b>17</b>	99.791377	93.593315	3	33.4
<b>18</b>	99.791667	95.798319	4	33.4
<b>19</b>	100.000000	89.859155	5	33.4
<b>20</b>	99.930216	91.483516	1	44.5
<b>21</b>	99.860627	87.292818	2	44.5
<b>22</b>	99.791377	93.593315	3	44.5
<b>23</b>	99.791667	95.238095	4	44.5
<b>24</b>	100.000000	89.859155	5	44.5
<b>25</b>	99.930216	91.483516	1	55.6
<b>26</b>	99.860627	86.740331	2	55.6
<b>27</b>	99.791377	93.593315	3	55.6
<b>28</b>	99.791667	95.238095	4	55.6
<b>29</b>	100.000000	89.859155	5	55.6
<b>30</b>	99.930216	91.483516	1	66.7
<b>31</b>	99.860627	86.740331	2	66.7


	<b>train</b>	<b>test</b>	<b>split</b>	<b>c</b>
<b>32</b>	99.791377	93.593315	3	66.7
<b>33</b>	99.791667	95.518207	4	66.7
<b>34</b>	100.000000	89.577465	5	66.7
<b>35</b>	99.930216	91.208791	1	77.8
<b>36</b>	99.860627	86.464088	2	77.8
<b>37</b>	99.791377	93.593315	3	77.8
<b>38</b>	99.791667	95.518207	4	77.8
<b>39</b>	100.000000	89.577465	5	77.8
<b>40</b>	99.930216	90.384615	1	88.9
<b>41</b>	99.860627	86.464088	2	88.9
<b>42</b>	99.791377	93.593315	3	88.9
<b>43</b>	99.791667	95.518207	4	88.9
<b>44</b>	100.000000	89.577465	5	88.9
<b>45</b>	99.930216	90.109890	1	100.0
<b>46</b>	99.860627	86.187845	2	100.0
<b>47</b>	99.791377	93.314763	3	100.0
<b>48</b>	99.791667	95.518207	4	100.0
<b>49</b>	100.000000	89.577465	5	100.0

```
In [93]: scores_l2.loc[scores_l2['split']==1]

for i in range(1,6):
    # p_train = pd.Series(scores_l2.loc[scores_l2['split']==i]["train"].values,
    # index=scores_l2.loc[scores_l2['split']==1]["c"])
    # # plt.figure()
    # p_train.plot()

    p_test = pd.Series(scores_l2.loc[scores_l2['split']==i]["test"].values, index=scores_l2.loc[scores_l2['split']==1]["c"])
    # plt.figure()
    p_test.plot()

plt.show()
```



```
In [99]: for c in c_vals:
    print(scores_l2.loc[scores_l2["c"]==c]["test"].mean(),c)
```

```
92.43629695082134 0.1
91.77242154289213 11.2
91.60357349252527 22.3
91.66036969166791 33.4
91.49337981879567 44.5
91.38288258122662 55.6
91.3825669620212 66.7
91.27237328829162 77.8
91.10753812345645 88.9
90.94163414332017 100.0
```

```
In [114]: # The highest C value is from 0.1, so using that to run on the data set

logreg_ridge_final = LogisticRegression(C=0.1)
x_train, x_test, y_train, y_test = train_test_split(digits.data, digits.target
, test_size=0.20, random_state=0)

logreg_ridge_final.fit(x_train, y_train) # fit
y_pred = logreg_ridge_final.predict(x_train) # predict
acc_log = sum(y_pred == y_train)/len(y_train)*100
print('Logistic Regression accuracy of train data:', str(round(acc_log,2)), '%')
)
print(confusion_matrix(y_train, y_pred))

y_pred = logreg_ridge_final.predict(x_test) # predict
acc_log = sum(y_pred == y_test)/len(y_test)*100
print('Logistic Regression accuracy of test data:', str(round(acc_log,2)), '%')
print(confusion_matrix(y_test, y_pred))
```

Logistic Regression accuracy of train data: 99.03 %

```
[[151  0  0  0  0  0  0  0  0  0]
 [  0 146  0  0  0  0  0  0  0  1]
 [  0  0 141  0  0  0  0  0  0  0]
 [  0  0  0 153  0  0  0  0  1  0]
 [  0  1  0  0 149  0  0  0  1  0]
 [  0  0  0  0  0 142  0  0  0  0]
 [  0  0  0  0  0  0 137  0  0  0]
 [  0  0  0  0  0  0  0 139  0  1]
 [  0  5  0  1  0  0  0  0 129  0]
 [  0  0  0  1  0  0  0  0  2 136]]
```

Logistic Regression accuracy of test data: 95.83 %

```
[[27  0  0  0  0  0  0  0  0  0]
 [  0 33  0  0  0  0  1  0  1  0]
 [  0  0 35  1  0  0  0  0  0  0]
 [  0  0  0 29  0  0  0  0  0  0]
 [  0  0  0  0 30  0  0  0  0  0]
 [  0  0  0  0  0 39  0  0  0  1]
 [  0  1  0  0  0  0 43  0  0  0]
 [  0  1  0  0  1  0  0 37  0  0]
 [  0  2  1  0  0  0  0  0 36  0]
 [  0  0  0  1  0  1  0  0  3 36]]
```

```
In [105]: # LASSO REGRESSION

c_vals_l1 = np.linspace(0.1,1000,10)

scores_l1_train = []
scores_l1_test = []
scores_l1_c = []
scores_l1_split = []

for c in c_vals_l1:
    # print("L1 Logistic Regression (Lasso)\n")
    split_count = 0
    logreg_lasso = LogisticRegression(penalty="l1", C=c)
    for train_index, test_index in skf.split(x, y):
        # print("TRAIN:", train_index, "TEST:", test_index)
        x_train, x_test = x[train_index], x[test_index]
        y_train, y_test = y[train_index], y[test_index]

        logreg_lasso.fit(x_train, y_train) # fit
        y_pred = logreg_lasso.predict(x_train) # predict
        acc_log = sum(y_pred == y_train)/len(y_train)*100
        # print('Logistic Regression accuracy of train data:', str(round(acc_
        _Log,2)), '%')
        # print(confusion_matrix(y_train, y_pred))
        scores_l1_train.append(acc_log)

        y_pred = logreg_lasso.predict(x_test) # predict
        acc_log = sum(y_pred == y_test)/len(y_test)*100
        # print('Logistic Regression accuracy of test data:', str(round(acc_
        Log,2)), '%')
        # print(confusion_matrix(y_test, y_pred))

    # print("\n")
    scores_l1_test.append(acc_log)
    scores_l1_c.append(c)
    split_count = split_count + 1
    scores_l1_split.append(split_count)
```



```
In [106]: data = {"train": scores_l1_train, "test": scores_l1_test, "c": scores_l1_c, "split": scores_l1_split}

scores_l1 = pd.DataFrame(data=data)

scores_l1[scores_l1.columns[::-1]]
```

Out[106]:

	<b>train</b>	<b>test</b>	<b>split</b>	<b>c</b>
<b>0</b>	98.534543	92.857143	1	0.1
<b>1</b>	98.606272	90.331492	2	0.1
<b>2</b>	98.331015	94.986072	3	0.1
<b>3</b>	98.263889	96.638655	4	0.1
<b>4</b>	98.474341	91.267606	5	0.1
<b>5</b>	99.930216	90.934066	1	111.2
<b>6</b>	99.930314	86.464088	2	111.2
<b>7</b>	99.930459	93.036212	3	111.2
<b>8</b>	99.930556	95.518207	4	111.2
<b>9</b>	100.000000	90.422535	5	111.2
<b>10</b>	99.930216	91.208791	1	222.3
<b>11</b>	100.000000	86.464088	2	222.3
<b>12</b>	99.930459	93.593315	3	222.3
<b>13</b>	99.930556	95.798319	4	222.3
<b>14</b>	100.000000	90.422535	5	222.3
<b>15</b>	99.930216	91.208791	1	333.4
<b>16</b>	100.000000	86.187845	2	333.4
<b>17</b>	99.930459	93.314763	3	333.4
<b>18</b>	99.930556	95.798319	4	333.4
<b>19</b>	100.000000	90.140845	5	333.4
<b>20</b>	99.930216	91.483516	1	444.5
<b>21</b>	99.860627	86.740331	2	444.5
<b>22</b>	99.930459	93.036212	3	444.5
<b>23</b>	100.000000	95.518207	4	444.5
<b>24</b>	100.000000	90.422535	5	444.5
<b>25</b>	99.930216	91.208791	1	555.6
<b>26</b>	99.930314	86.464088	2	555.6
<b>27</b>	99.930459	93.036212	3	555.6
<b>28</b>	100.000000	95.518207	4	555.6
<b>29</b>	100.000000	90.422535	5	555.6
<b>30</b>	99.930216	91.208791	1	666.7
<b>31</b>	100.000000	85.911602	2	666.7


	<b>train</b>	<b>test</b>	<b>split</b>	<b>c</b>
<b>32</b>	99.930459	92.757660	3	666.7
<b>33</b>	100.000000	95.518207	4	666.7
<b>34</b>	100.000000	90.422535	5	666.7
<b>35</b>	99.930216	91.208791	1	777.8
<b>36</b>	100.000000	85.911602	2	777.8
<b>37</b>	99.930459	93.036212	3	777.8
<b>38</b>	100.000000	95.798319	4	777.8
<b>39</b>	100.000000	90.704225	5	777.8
<b>40</b>	99.930216	91.208791	1	888.9
<b>41</b>	100.000000	86.187845	2	888.9
<b>42</b>	99.930459	93.036212	3	888.9
<b>43</b>	100.000000	95.798319	4	888.9
<b>44</b>	100.000000	90.422535	5	888.9
<b>45</b>	99.930216	90.934066	1	1000.0
<b>46</b>	100.000000	86.464088	2	1000.0
<b>47</b>	99.930459	92.757660	3	1000.0
<b>48</b>	99.930556	95.518207	4	1000.0
<b>49</b>	100.000000	90.422535	5	1000.0

```
In [109]: # scores_l1.loc[scores_l1['split']==1]

for i in range(1,6):
#     p_train = pd.Series(scores_l2.loc[scores_l2['split']==i]["train"].values, index=scores_l2.loc[scores_l2['split']==1]["c"])
#     plt.figure()
#     p_train.plot()

    p_test = pd.Series(scores_l1.loc[scores_l1['split']==i]["test"].values, index=scores_l1.loc[scores_l1['split']==1]["c"])
#     plt.figure()
    p_test.plot()

plt.show()
```



```
In [111]: for c in c_vals_l1:
            print(scores_l1.loc[scores_l1["c"]==c]["test"].mean(),c)
```

```
93.21619361784721 0.1
91.27502170504022 111.2
91.49740978176223 222.3
91.33011282840201 333.4
91.44016043371487 444.5
91.32996675998528 555.6
91.16375921600952 666.7
91.3318299595488 777.8
91.33074055016434 888.9
91.21931139863354 1000.0
```

```
In [115]: # The highest C value is from 0.1, so using that to run on the data set

logreg_lasso_final = LogisticRegression(penalty="l1", C=0.1)
x_train, x_test, y_train, y_test = train_test_split(digits.data, digits.target
, test_size=0.20, random_state=0)

logreg_lasso_final.fit(x_train, y_train) # fit
y_pred = logreg_lasso_final.predict(x_train) # predict
acc_log = sum(y_pred == y_train)/len(y_train)*100
print('Logistic Regression accuracy of train data:', str(round(acc_log,2)), '%')
)
print(confusion_matrix(y_train, y_pred))

y_pred = logreg_lasso_final.predict(x_test) # predict
acc_log = sum(y_pred == y_test)/len(y_test)*100
print('Logistic Regression accuracy of test data:', str(round(acc_log,2)), '%')
print(confusion_matrix(y_test, y_pred))
```

Logistic Regression accuracy of train data: 98.05 %

```
[[151  0  0  0  0  0  0  0  0  0]
 [  0 145  0  0  0  0  0  0  1  1]
 [  0  0 141  0  0  0  0  0  0  0]
 [  0  0  0 149  0  0  0  1  3  1]
 [  0  2  0  0 148  0  0  0  1  0]
 [  0  0  0  0  1 140  1  0  0  0]
 [  0  0  0  0  0  1 135  0  1  0]
 [  0  0  0  0  0  0  0 138  1  1]
 [  0  6  0  1  0  1  0  0 127  0]
 [  0  0  0  1  0  1  0  0  2 135]]
```

Logistic Regression accuracy of test data: 95.83 %

```
[[27  0  0  0  0  0  0  0  0  0]
 [  0 33  0  0  0  0  1  0  1  0]
 [  1  0 35  0  0  0  0  0  0  0]
 [  0  0  0 29  0  0  0  0  0  0]
 [  0  0  0  0 29  0  0  1  0  0]
 [  0  0  0  0  0 39  0  0  0  1]
 [  0  0  0  0  0  0 44  0  0  0]
 [  0  0  0  0  1  0  0 38  0  0]
 [  0  2  1  0  0  0  0  0 36  0]
 [  0  1  0  0  0  1  0  1  3 35]]
```